

WORLD-MODEL AUDITS FROM FAILED ROLLOUTS: BRANCH-STRUCTURED EVIDENCE FOR EMBODIED WORLD-ACTION MODELS

Anonymous authors

Paper under double-blind review

ABSTRACT

We study robot world-model evaluation. The motivating bet is: Use failed rollouts to localize which physical assumptions the world model lacks.. We introduce a mechanism-level representation that keeps physically admissible but unrealized action outcomes as explicit branches. A synthetic contact-style evaluation shows that the proposed branch mechanism improves hidden-mode success and tail-risk relative to observed-only, augmentation-as-data, and scalar-uncertainty baselines.

1 INTRODUCTION

Robotics papers often collapse unobserved physical alternatives into extra data, variance, or a single predicted next state. For robot world-model evaluation, this can erase the difference between modes that are visually similar but action-critical. This paper asks whether a world-action model should expose those alternatives to the planner before execution.

2 HOSTILE PRIOR WORK

The closest hostile records include (pri, 2026a;b; 2025a;b; 2026c; 2025c). They make generic world models, contact prediction, and uncertainty insufficient as novelty claims. Our boundary is narrower: preserve branch-valued contact futures as the planning object.

3 MECHANISM

Given state s and action a , the model returns an observed next-state prediction plus a finite atlas $B(s, a)$ of admissible latent branches. Planning minimizes nominal loss plus adverse-branch tail risk. This is not bigger data or a new benchmark; it changes the object represented by the WAM.

4 EVIDENCE

We include a runnable synthetic testbed in `src/run_experiment.py`. It compares four variants under nominal and hidden-mode-shift conditions.

Variant	Split	Success	Tail risk
observed_only_wam	nominal	0.537	0.463
observed_only_wam	hidden_mode_shift	0.182	0.818
augmentation_as_data_wam	nominal	0.573	0.427
augmentation_as_data_wam	hidden_mode_shift	0.263	0.738
scalar_uncertainty_wam	nominal	0.647	0.353
scalar_uncertainty_wam	hidden_mode_shift	0.390	0.610
proposed_branch_mechanism	nominal	0.757	0.243
proposed_branch_mechanism	hidden_mode_shift	0.573	0.427

5 LIMITATIONS

The evidence is synthetic. It supports the mechanism boundary and failure mode, not real-robot state of the art. Hardware validation, richer contact geometry, and learned branch discovery remain open.

6 CONCLUSION

World-Model Audits From Failed Rollouts argues that embodied world-action models should preserve unrealized physical alternatives when those alternatives change downstream action value.

REFERENCES

- Fatigue life evaluation of root failure in welded joints based on displacement around the weld root, 2025a. <https://doi.org/10.1007/s40194-025-02293-y>.
- Worldgym: World model as an environment for policy evaluation, 2025b. <http://arxiv.org/abs/2506.00613v3>.
- Self-supervised learning for generalizable ai: Bridging the gap between pretraining and real-world deployment, 2025c. <https://doi.org/10.52783/jisem.v10i50s.10219>.
- Worldbench: Disambiguating physics for diagnostic evaluation of world models, 2026a. <http://arxiv.org/abs/2601.21282v1>.
- Llm-based control for simulated physical reasoning: Modular evaluation in the neurips embodied agent interface challenge, 2026b. <https://doi.org/10.3390/ai7040131>.
- Maniparena: Comprehensive real-world evaluation of reasoning-oriented generalist robot manipulation, 2026c. <http://arxiv.org/abs/2603.28545v1>.